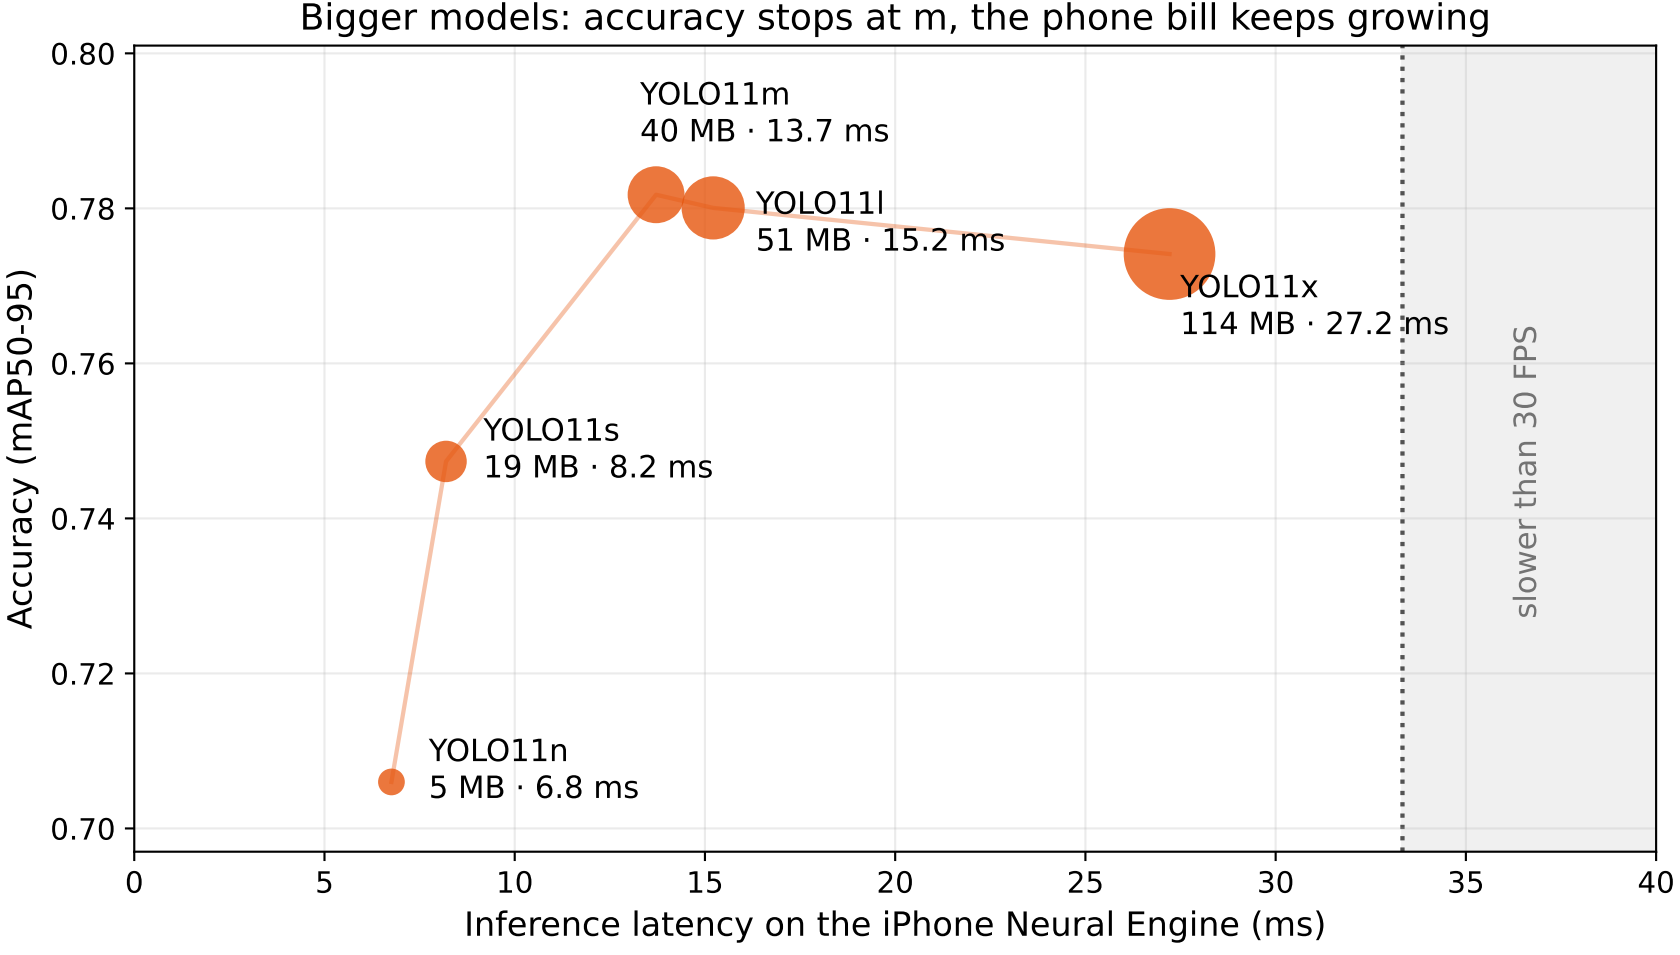
Bigger models: accuracy stops at m, the phone bill keeps growing
Accuracy (mAP50-95)
Inference latency on the iPhone Neural Engine (ms)
0.80
0.78
0.76
0.74
0.72
0.70
0
5
10
15
20
25
30
35
40
YOLO11m
40 MB · 13.7 ms
YOLO11l
51 MB · 15.2 ms
YOLO11x
114 MB · 27.2 ms
YOLO11s
19 MB · 8.2 ms
YOLO11n
5 MB · 6.8 ms
slower than 30 FPS